

Agentic PMO Suite

Named AI agents that draft every standard PMO artifact, triage intake, and score prioritization — grounded in real program context, with humans always in control.

STATUS: WORK IN PROGRESS — PROTOTYPE, NOT A FINISHED PRODUCT ~75–80% of its own scoped goal; ~35–40% of a full enterprise PfMO function. Full breakdown in the app's own Roadmap & Scope tab.

Why: PMOs run on the same handful of artifacts — charters, plans, status reports, risk registers — rewritten by hand for every program, and prioritization decisions that are hard to defend when leadership asks “why this and not that?” This suite uses a set of purpose-built AI agents, not one generic chatbot, to draft each artifact consistently and to score prioritization against a fixed, OKR-gated rubric — while keeping every override badged and reversible back to the agent's original assessment.

Capability	What It Does	Why It Matters
Intake & AI Triage	A short submission form (name, description, sponsor, rough budget) that an Intake Agent triages: drafts the problem statement, checks value-prop quality, assesses OKR alignment, and drafts requirements/success criteria/scope/risks for human review.	Keeps submission fast while producing PMO-grade write-ups — no request skips scrutiny, and no submitter is blocked by a long form.
OKR-Gated Prioritization Scoring	A Prioritization Agent scores every request against a fixed rubric: OKR tie (hard gate — no tie, no score), size, team spread/global reach, complexity, and a 9-point value component.	Answers leadership's “are we working on the right things?” with one consistent, auditable rubric instead of the loudest voice in the room.
Product vs. Platform Value Tracks	Product work is scored on customer impact, new customers, and net-new revenue. Platform/infra work is scored on existing-customer coverage, churn risk, and revenue at risk — same 9-point scale, mirrored structure.	Lets infrastructure and reliability work compete fairly for resourcing instead of losing every time to initiatives with a revenue number attached.
Human Override, Badged & Reversible	Any agent score or drafted content can be overridden by a PMO or engineering leader. Overrides are clearly badged and always revertible back to the agent's original assessment. (A timestamped who/when/why audit log is roadmapped, not yet built.)	Keeps a human accountable for every resourcing call, without losing the agent's original reasoning or silently overwriting it.
TPM Capacity Governance	No Technical Program Manager can be allocated above 80% of stated capacity — enforced automatically on intake approval and on manual allocation changes.	Prevents the quiet over-commitment of scarce TPM bandwidth that shows up as missed dates three months later.
Seven Grounded Artifact Agents	Charter, Project Plan, Status Report, Risk Register, Dependency Tracker, OKR Report, and Program Landing Page — each drafted from one shared, real program-context object and a fixed section template.	Every artifact says the same thing about the same program. No drift between what the charter promises and what the status report reports.
Risk Register with Organizational Memory	A seeded Risk Pattern Library of risks tagged by initiative type, teams, size, and domain keywords is matched against each new program and surfaced before generation; the Risk Register Agent is told to include matched patterns first, adapted to the specifics, then add net-new risks.	Recurring risks (vendor-change breaks, cutover data drift, API compatibility gaps) get caught because the system remembers them — not because an agent happens to reinvent them correctly each time.
Voice / Transcript Intake	A requester can paste a meeting transcript (or use best-effort live dictation) and a Transcript Parser Agent fills the intake form fields directly.	Removes the “yet another form” friction that quietly degrades request quality and discourages submission altogether.
OKR / Metric Consistency Check	A structural, rule-based check (not an agent) flags when a program's named OKR doesn't actually match the category of metric it's tracking (e.g., a cost-reduction OKR paired with churn-risk figures).	Catches a real data-quality gap systematically and every time, rather than hoping an agent happens to notice it.
Consolidated Program Tracker	One portfolio view of every active program — RAG status, phase, budget, TPM allocation, and prioritization score — alongside full detail in each program's own tab.	One place to see the whole portfolio instead of reconciling five different decks before a leadership review.

Architecture: a single self-contained web app; each “agent” is a distinctly-prompted call to the same model, separated by prompt design and a structured output contract rather than by infrastructure. Code: github.com/VjRa03An/automation-agents/tree/main/agentic-pmo-suite